

Machine Learning Assignment — Guidelines

Objective

The objective of this assignment is to enable students to collect a local dataset, apply a new machine learning algorithm, evaluate and explain the model, and optionally integrate it into a front-end system.

Task Description

Students must identify a real-world problem, collect or compile a dataset, apply a machine learning algorithm not taught in lectures, train and evaluate it, and explain the results using XAI techniques. **Avoid developing an image processing application.**

Guidelines

1. Problem Definition & Dataset Collection (**15 marks**)

- Clearly describe the problem and its relevance.
- Explain:
 - Data source (how and from where it was collected),
 - Features and target variable,
 - Size of the dataset,
- Any preprocessing done (cleaning, encoding, normalization).
 - Ensure ethical data use (no personal or sensitive data without consent).

2. Selection of a New Machine Learning Algorithm (**15 marks**)

- Choose an algorithm **not covered during lectures** and avoid using deep learning models.
- Justify:
 - Why was this algorithm selected,
 - How it differs from standard models (e.g., decision trees, logistic regression, k-NN, etc.).

3. Model Training and Evaluation (20 marks)

- Explain:
 - Train/validation/test split,
 - Hyperparameter choices,
 - Performance metrics used (accuracy, F1, RMSE, AUC, etc. depending on task),
 - Results obtained and what they indicate.
- Include tables, graphs, or plots where appropriate.

4. Explainability & Interpretation (20 marks)
 - Apply at least **one explainability method**, such as:
 - SHAP
 - LIME
 - Feature importance analysis
 - Partial Dependence Plots (PDP)
 - Explain:
 - What the model has learned,
 - Which features are most influential,
 - Whether the model's behavior aligns with domain knowledge.
5. Critical Discussion (10 marks)
 - Limitations of the model,
 - Data quality issues,
 - Risks of bias or unfairness,
 - Potential real-world impact and ethical considerations.
6. Report Quality & Technical Clarity (10 marks)
7. Bonus: Front-End Integration (10 marks)

Bonus marks will be awarded for:

 - Integrating the trained model into a front-end system (web app, dashboard, mobile app, etc.).
 - Allowing users to input data and view predictions/explanations.
 - Examples: Streamlit app, Flask + HTML, React frontend, etc.

Submission Requirements

- **Written Report (PDF)**

Including problem description, methodology, results, interpretation, and discussion.

- **Source Code (ZIP / GitHub link)**

Including data preprocessing, training, evaluation, and explainability scripts.

- **Dataset (if publicly shareable)** or a description of how it was obtained.

- **Demo video (3–5 minutes)** showing the front-end system.

Upload your submission to the Moodle course page.